

Adversarial Variational Bayes in Edward

Karttikeya Mangalam, 14311

Shubh Gupta, 14670

Ankur Kumar, 14109

Course Project

CS698X - Topics in Probabilistic Modeling and Inference

Prof. Piyush Rai

29 avril 2018



- 1 Introduction
- 2 Proposed Approach
- 3 Algorithm
- 4 Results

Overview

- Generative Adversarial Networks & Variational Auto-encoders are leading generative models both of which are probabilistic in nature.
- GANs produce sharper image with better visual quality
- VAEs produce a blurry images but have attractive mathematical bounds



FIGURE – Visual output from VAEs¹

1. Tulyakov et al. ICLR 2017

Adversarial Variational Auto-encoders

- Attempts to combine best of both worlds
- Makzani et al. proposed Adversarial Auto-encoder (AAE)
 - ▶ Doesn't lead to MLE assignments in general
 - ▶ Don't optimize a lower bound on maximum likelihood
- Mescheder et al. propose an alternate approach : Adversarial VAE
 - ▶ Optimize ELBO lower bound to maximum likelihood
 - ▶ AAEs can be interpreted as an approximation
 - ▶ Uses 2-player game like GANs to circumvent evaluating likelihood

- 1 Introduction
- 2 Proposed Approach**
- 3 Algorithm
- 4 Results

Framework

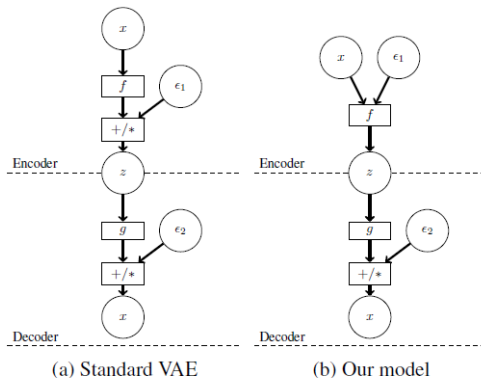


FIGURE – The vanilla VAE model and the modification proposed by Mescheder et al. VAEs are specified by a parametric generative model $p(x_j|z)$ of the visible variables given the latent variables, a prior $p(z)$ over the latent variables and an approximate inference model $q(z_j|x)$ over the latent variables given the visible variables.

ELBO Formulation I

- Standard Evidence Lower bound for Variational Auto-encoders

$$\log p_{\theta}(x) \geq -\text{KL}(q_{\phi}(z | x), p(z)) \\ + \mathbb{E}_{q_{\phi}(z|x)} \log p_{\theta}(x | z)$$

- Our Goal is to then solve the following optimization problem

$$\max_{\theta} \max_{\phi} \mathbb{E}_{p_{\mathcal{D}}(x)} \mathbb{E}_{q_{\phi}(z|x)} (\log p(z) \\ - \log q_{\phi}(z | x) + \log p_{\theta}(x | z)).$$

ELBO Formulation II

- Rewriting ELBO as

$$\max_{\theta} \max_{\phi} \mathbb{E}_{p_{\mathcal{D}}(x)} \mathbb{E}_{q_{\phi}(z|x)} (\underbrace{\log p(z)}_{\underbrace{-\log q_{\phi}(z|x)} + \log p_{\theta}(x|z)}).$$

- The indicated terms can be re-written as the optimal value of an additional real-valued discriminative network $T(x; z)$ defined as

$$\begin{aligned} \max_T \mathbb{E}_{p_{\mathcal{D}}(x)} \mathbb{E}_{q_{\phi}(z|x)} \log \sigma(T(x, z)) \\ + \mathbb{E}_{p_{\mathcal{D}}(x)} \mathbb{E}_{p(z)} \log (1 - \sigma(T(x, z))) \end{aligned}$$

Can be thought of as a two player min max game similar to GANs.

ELBO Formulation III

Thus we have the optimal value $T^*(x; z)$ is obtained as

$$T^*(x, z) = \log q_\phi(z | x) - \log p(z)$$

$$\max_{\theta, \phi} \mathbb{E}_{p_{\mathcal{D}}(x)} \mathbb{E}_{q_\phi(z|x)} (-T^*(x, z) + \log p_\theta(x | z))$$

Now what remains is to find the gradients of the above loss with respect to θ and ϕ that can be used in any first order gradient optimizer. Here we employ the result,

$$\mathbb{E}_{q_\phi(z|x)} (\nabla_\phi T^*(x, z)) = 0.$$

And taking gradient for θ is straightforward.

- 1 Introduction
- 2 Proposed Approach
- 3 Algorithm**
- 4 Results

Algorithm

Now, reparameterization trick is used to rewrite the ELBO loss as

$$\max_{\theta, \phi} \mathbb{E}_{p_{\mathcal{D}}(x)} \mathbb{E}_{\epsilon} \left(-T^*(x, z_{\phi}(x, \epsilon)) \right. \\ \left. + \log p_{\theta}(x \mid z_{\phi}(x, \epsilon)) \right)$$

Thus, the parameters can be updated as :

$$\begin{aligned} g_{\theta} &\leftarrow \frac{1}{m} \sum_{k=1}^m \nabla_{\theta} \log p_{\theta} \left(x^{(k)} \mid z_{\phi} \left(x^{(k)}, \epsilon^{(k)} \right) \right) \\ g_{\phi} &\leftarrow \frac{1}{m} \sum_{k=1}^m \nabla_{\phi} \left[-T_{\psi} \left(x^{(k)}, z_{\phi} \left(x^{(k)}, \epsilon^{(k)} \right) \right) \right. \\ &\quad \left. + \log p_{\theta} \left(x^{(k)} \mid z_{\phi} \left(x^{(k)}, \epsilon^{(k)} \right) \right) \right] \\ g_{\psi} &\leftarrow \frac{1}{m} \sum_{k=1}^m \nabla_{\psi} \left[\log \left(\sigma \left(T_{\psi} \left(x^{(k)}, z_{\phi} \left(x^{(k)}, \epsilon^{(k)} \right) \right) \right) \right) \right. \\ &\quad \left. + \log \left(1 - \sigma \left(T_{\psi} \left(x^{(k)}, z_{\phi} \left(x^{(k)}, \epsilon^{(k)} \right) \right) \right) \right) \right] \end{aligned}$$

Implementation Details

- Prior over latent variable z is Normal Distribution and the generated data x fits a Bernoulli distribution with output of generative network as logits.

Implementation Details

- Prior over latent variable z is Normal Distribution and the generated data x fits a Bernoulli distribution with output of generative network as logits.
- Inference network is modeled by adding Gaussian noise ϵ to observed data.

Implementation Details

- Prior over latent variable z is Normal Distribution and the generated data x fits a Bernoulli distribution with output of generative network as logits.
- Inference network is modeled by adding Gaussian noise ϵ to observed data.
- The parameters of variational distribution of z (a Normal distribution) is found by optimizing the previous ELBO formulation.

Implementation Details

- Prior over latent variable z is Normal Distribution and the generated data x fits a Bernoulli distribution with output of generative network as logits.
- Inference network is modeled by adding Gaussian noise ϵ to observed data.
- The parameters of variational distribution of z (a Normal distribution) is found by optimizing the previous ELBO formulation.
- Optimizing ELBO with respect to ϕ while keeping θ and T fixed makes the encoder network collapse to a deterministic f .

Implementation Details

- Prior over latent variable z is Normal Distribution and the generated data x fits a Bernoulli distribution with output of generative network as logits.
- Inference network is modeled by adding Gaussian noise ϵ to observed data.
- The parameters of variational distribution of z (a Normal distribution) is found by optimizing the previous ELBO formulation.
- Optimizing ELBO with respect to ϕ while keeping θ and T fixed makes the encoder network collapse to a deterministic f .
- Perform several SGD-updates for the adversary for one SGD-update of the generative network.

Implementation Details

- Prior over latent variable z is Normal Distribution and the generated data x fits a Bernoulli distribution with output of generative network as logits.
- Inference network is modeled by adding Gaussian noise ϵ to observed data.
- The parameters of variational distribution of z (a Normal distribution) is found by optimizing the previous ELBO formulation.
- Optimizing ELBO with respect to ϕ while keeping θ and T fixed makes the encoder network collapse to a deterministic f .
- Perform several SGD-updates for the adversary for one SGD-update of the generative network.
- Inference, Generator and Discriminator network uses fully connected layers.


Implementation Details

- Prior over latent variable z is Normal Distribution and the generated data x fits a Bernoulli distribution with output of generative network as logits.
- Inference network is modeled by adding Gaussian noise ϵ to observed data.
- The parameters of variational distribution of z (a Normal distribution) is found by optimizing the previous ELBO formulation.
- Optimizing ELBO with respect to ϕ while keeping θ and T fixed makes the encoder network collapse to a deterministic f .
- Perform several SGD-updates for the adversary for one SGD-update of the generative network.
- Inference, Generator and Discriminator network uses fully connected layers.
- ADAM optimizer for Discriminative network, SGD for Generative Network.

- 1 Introduction
- 2 Proposed Approach
- 3 Algorithm
- 4 Results**


Synthetic Dataset I

- Synthetic data generated as one-hot vectors of feature size 4, to be encoded into 2 dimensional latent variables.




Synthetic Dataset II

- Synthetic data generated as one-hot vectors of feature size 8 with 2 dimensional latent space visualization



Synthetic Dataset III

- Synthetic data generated as binarized features of feature size 8 with 2 dimensional latent space visualization



MNIST Experiment I

- Visualization of 2 dimensional embeddings over binarized MNIST dataset.





FIGURE – a) 500 iterations b) 2000 iterations c) 5000 iterations



MNIST Experiment II

- T-SNE visualization over 10 dimensional latent space embeddings of binarized MNIST



MNIST Experiment III

- Sampled digits from our implemented model on MNIST



CelebA Dataset

- Samples from results of original paper on CelebA dataset.



Log likelihoods

	$\log p(x) \geq$	$\log p(x) \approx$
AVB (8-dim)	$(\approx -83.6 \pm 0.4)$	-91.2 ± 0.6
AVB + AC (8-dim)	$\approx -96.3 \pm 0.4$	-89.6 ± 0.6
EDWARD AVB	$\approx -96.7 \pm 0.5$	-
VAE (8-dim)	-98.1 ± 0.5	-90.9 ± 0.6

TABLE – Log likelihood values for different approaches on binarized MNIST dataset.

- Our implementation is competitive although the slight loss in performance is due to using few fully connected layers instead of deep convolutional layers for discriminator network.

Conclusion

- Programming Probabilistic models is easier in Edward as compared to other non-probabilistic frameworks with lesser lines of code.
- Training time taken for Edward is more than other frameworks, mostly because of high level abstraction and hence less customization.
- GPU utilization for Edward programs is more than that for other probabilistic frameworks, since it uses Tensorflow's computation capabilities.
- Edward is still in development phase, and hence many models, such as LDA, are hard to train, even if modeling is easier.